

# Methods for detecting cyber attacks

## Introduction To ML Part 2

Dr. Ran Dubin

Lecture slides were taken from Prof. Lior Rokach  
and Asaf Shabtai

# Performance evaluation

*Evaluating what's been learned*



# Introduction

- How to measure algorithm success?
- Compare between algorithms
- Training/validation/test sets
- Linear Regression
- Resampling methods: k-fold cross-validation
- Feature selection

# Algorithm preference

- Criteria (application-dependent):
  - Accuracy
  - Misclassification error, or risk (loss functions)
  - Training time/space complexity
  - Testing time/space complexity
  - Interpretability (model complexity, for instance number of hidden units)
  - Easy tuning
  - Easy programmability
  - Easy embedding

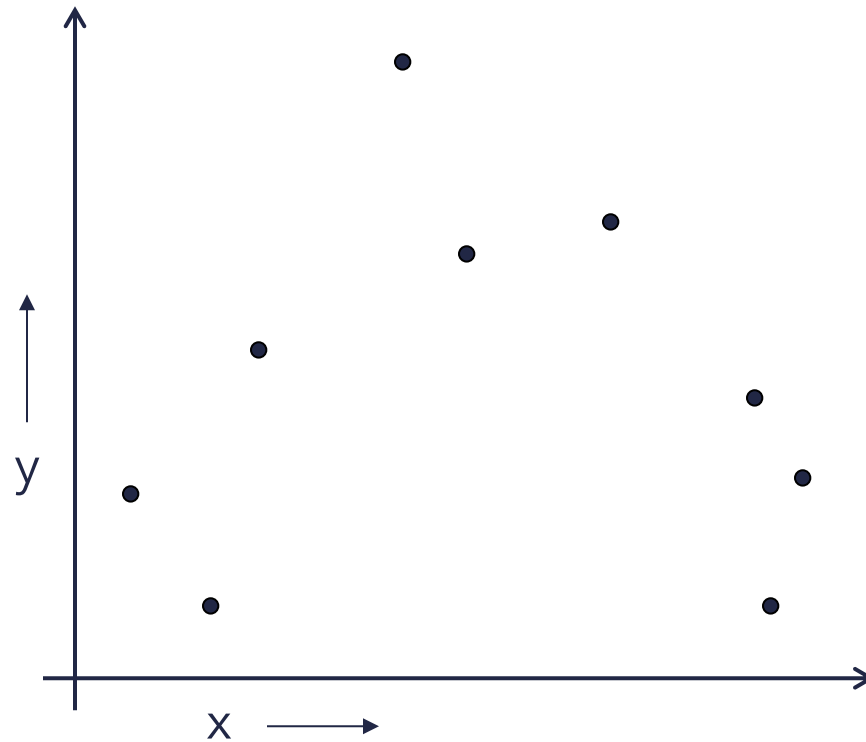


**"Fast is fine, but  
accuracy is everything."**  
Xenophon  
(Greek historian,  
**430 - 354 BC**)

# Other criteria

- Robustness
- Missing values
- Noise
- Incremental
- Data type (e.g. numeric)
- Multi-class
- Imbalance
- Cost sensitive
- Availability

# A Regression Problem

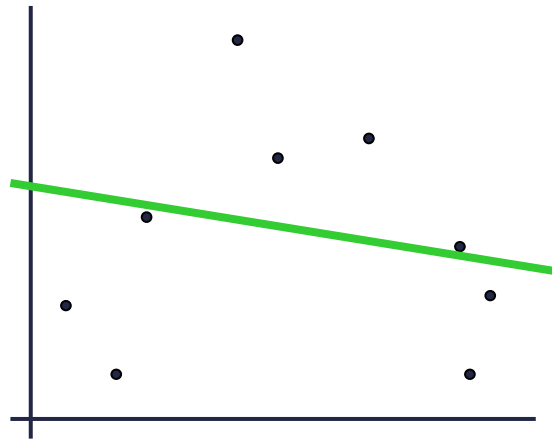


$$y = f(x) + \text{noise}$$

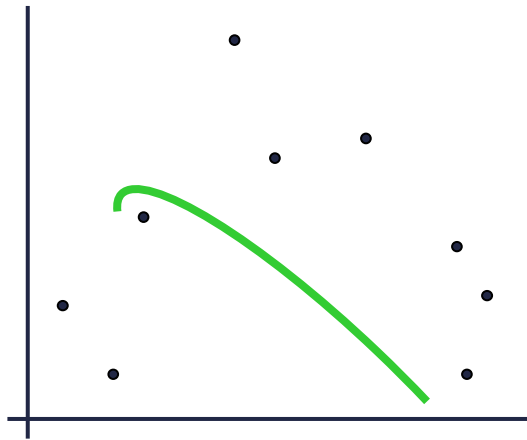
Can we learn  $f$  from this data?

Let's consider three methods...

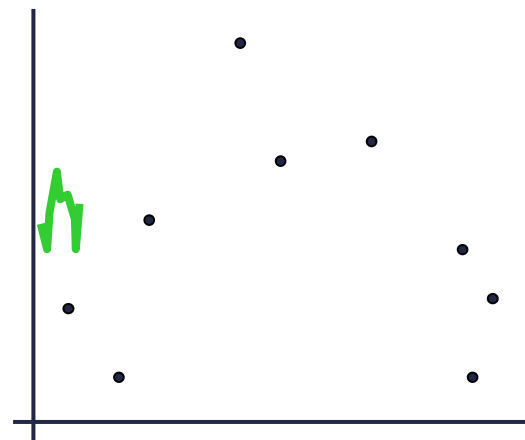
# Which is best?



*Linear regression*



*Quadratic regression*

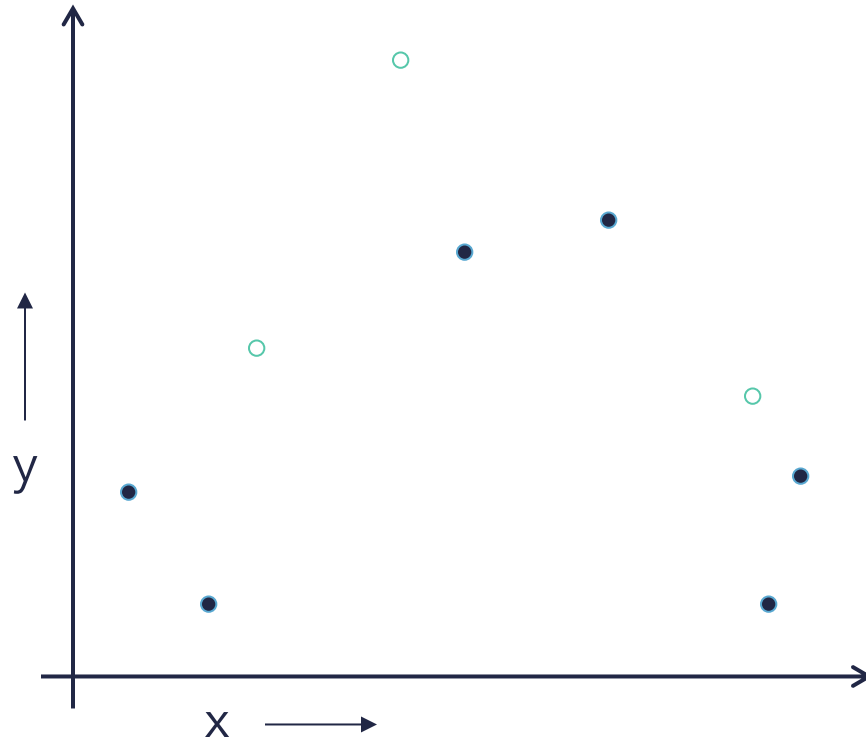


Piecewise linear  
nonparametric regression

Why not choose the method with  
the best fit to the data?

"How well are you going to  
predict future data drawn  
from the same distribution?"

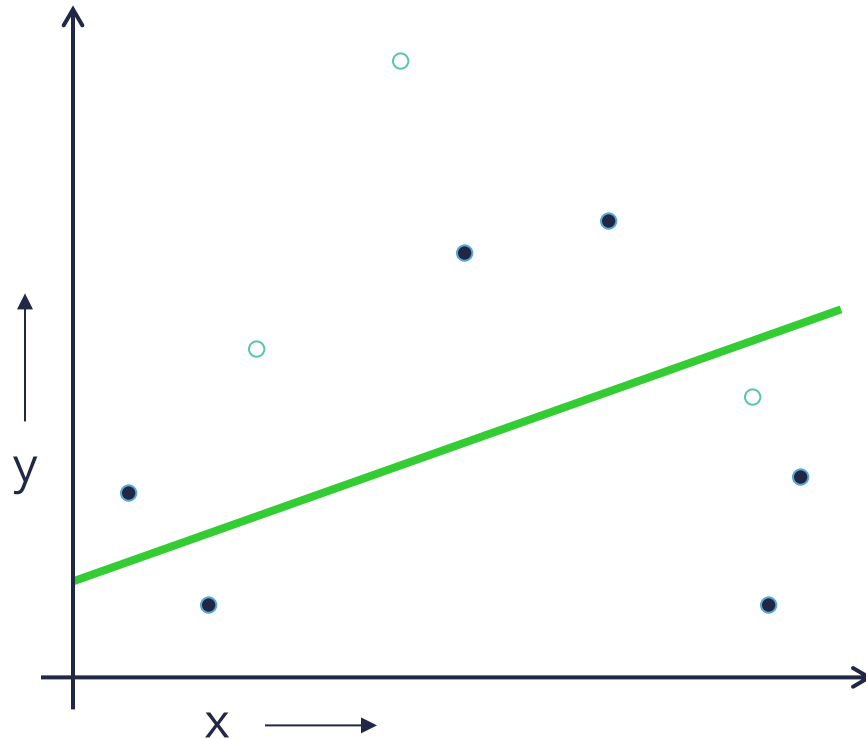
# The test set method



1. Randomly choose 30% of the data to be in a **test set**
2. The remainder is a **training set**



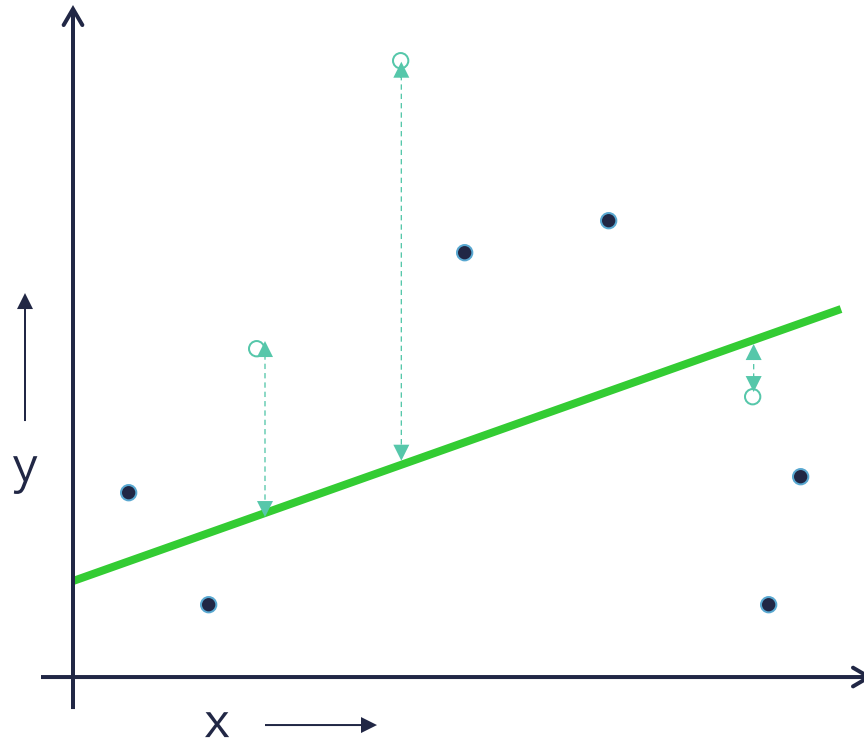
# The test set method



Linear regression example

1. Randomly choose 30% of the data to be in a **test set**
2. The remainder is a **training set**
3. Perform your regression on the training set

# The test set method



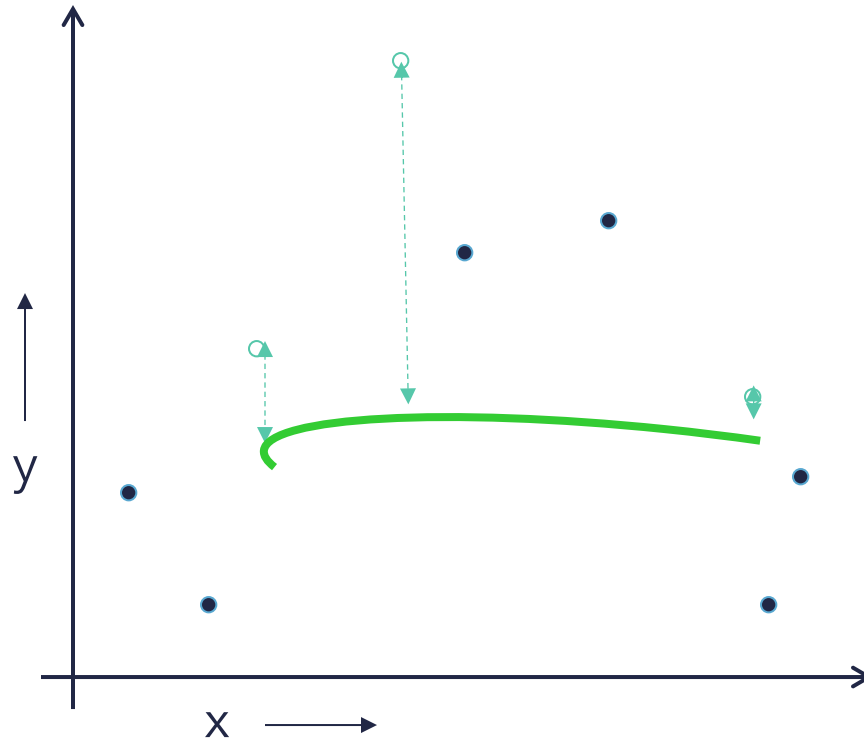
Linear regression example  
Mean Squared Error = 2.4

1. Randomly choose 30% of the data to be in a **test set**
2. The remainder is a **training set**
3. Perform your regression on the training set
4. Estimate your future performance with the test set

Note: MSE == 0 perfect regression

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (\hat{Y}_i - Y_i)^2 :$$

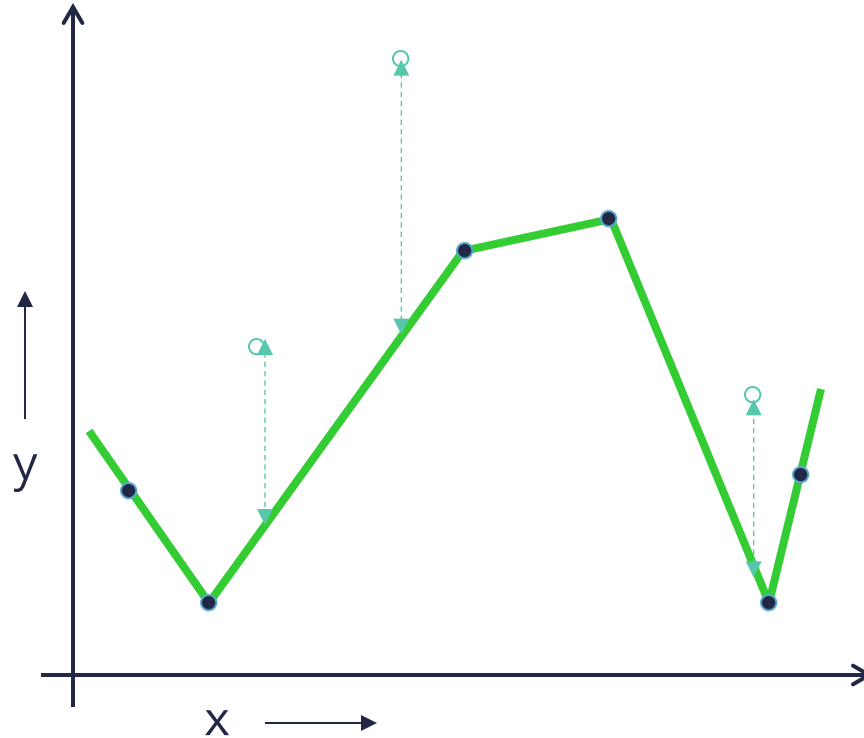
# The test set method



Quadratic regression example  
Mean Squared Error = 0.9

1. Randomly choose 30% of the data to be in a **test set**
2. The remainder is a **training set**
3. Perform your regression on the training set
4. Estimate your future performance with the test set

# The test set method



Join-the-dots example  
Mean Squared Error = 2.2

1. Randomly choose 30% of the data to be in a **test set**
2. The remainder is a **training set**
3. Perform your regression on the training set
4. Estimate your future performance with the test set

# The test set method

- Good news:
  - Simple
  - Can then simply choose the method with the best test-set score
- Bad news:
  - What's the downside?

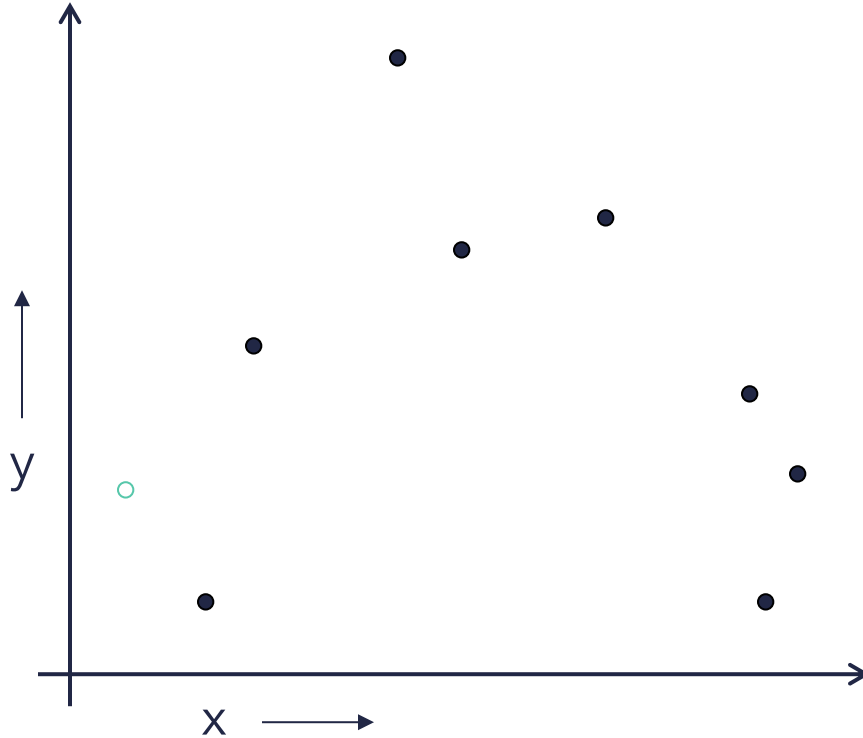
# The test set method

- Good news:
  - Simple
  - Can then simply choose the method with the best test-set score
- Bad news:
  - Wastes data: we get an estimate of the best method to apply to 30% less data
  - If we don't have much data, our test-set might just be lucky or unlucky
  - The "test-set estimator of performance has high variance"

# LOOCV (Leave-one-out Cross Validation)

For  $k=1$  to  $R$

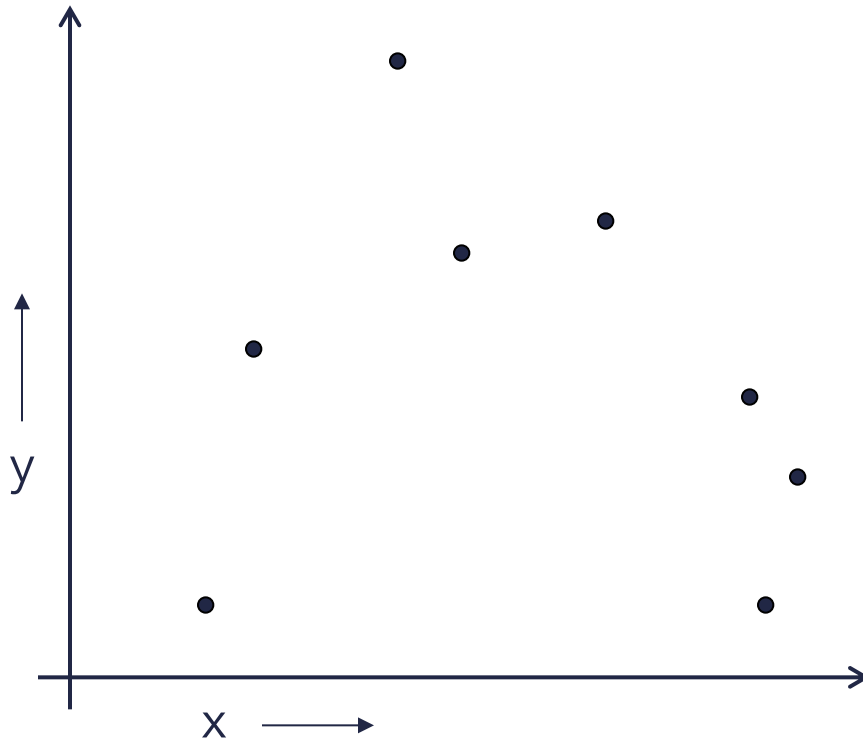
1. Let  $(x_k, y_k)$  be the  $k^{\text{th}}$  record



# LOOCV (Leave-one-out Cross Validation)

For  $k=1$  to  $R$

1. Let  $(x_k, y_k)$  be the  $k^{\text{th}}$  record
2. Temporarily remove  $(x_k, y_k)$  from the dataset

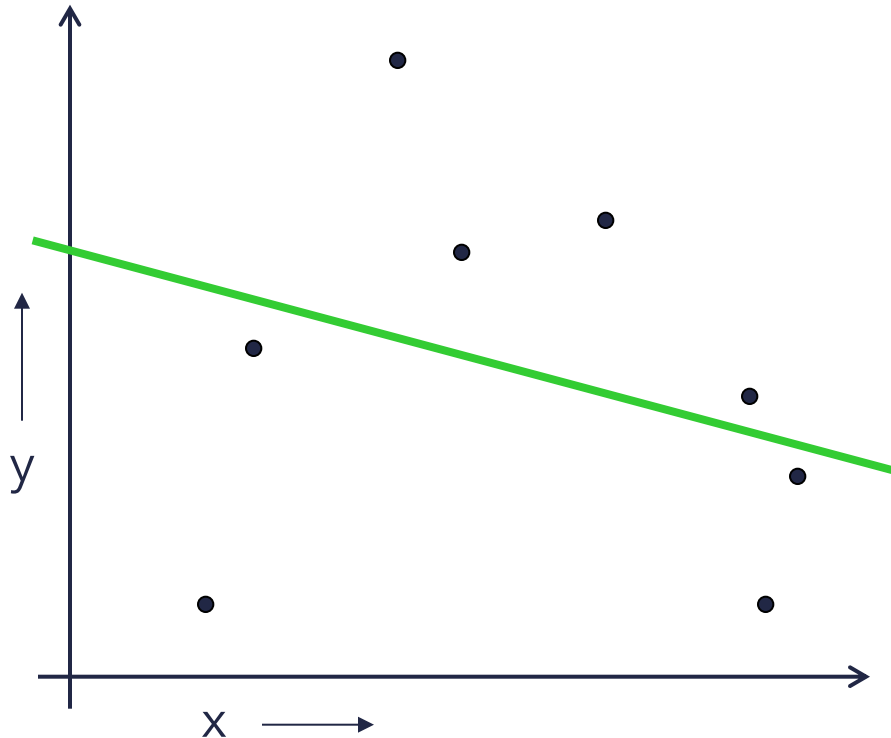




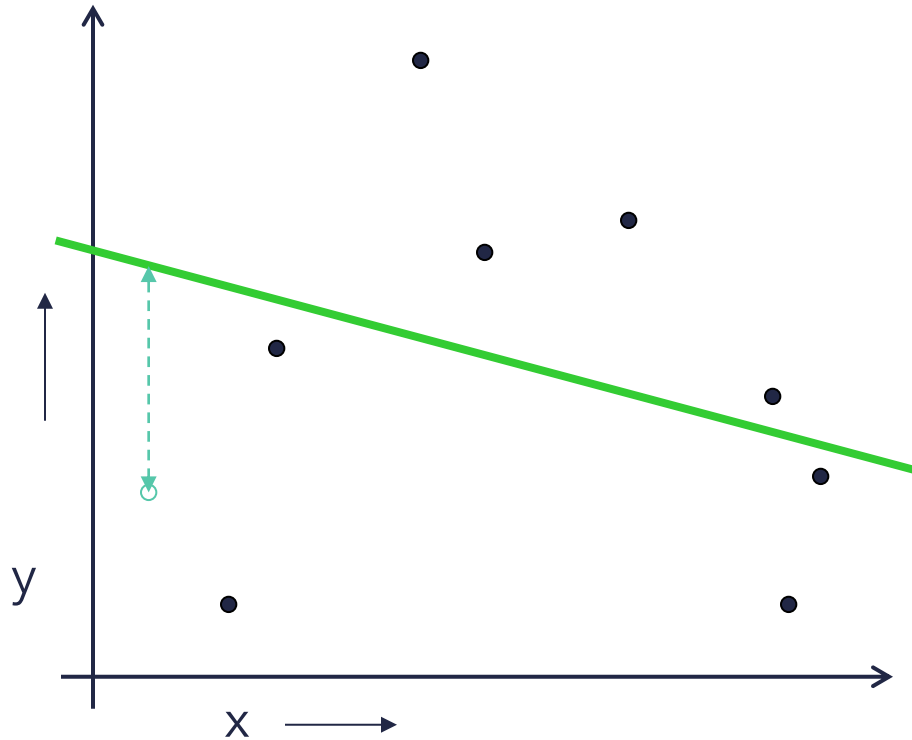
# LOOCV (Leave-one-out Cross Validation)

For  $k=1$  to  $R$

1. Let  $(x_k, y_k)$  be the  $k^{\text{th}}$  record
2. Temporarily remove  $(x_k, y_k)$  from the dataset
3. Train on the remaining  $(R-1)$  datapoints



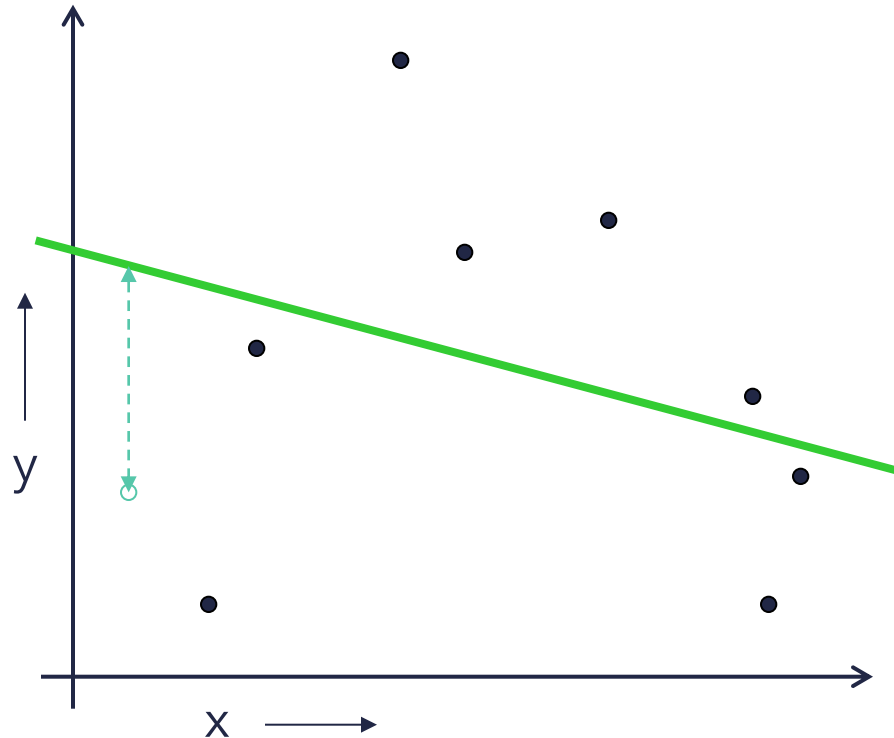
# LOOCV (Leave-one-out Cross Validation)



For  $k=1$  to  $R$

1. Let  $(x_k, y_k)$  be the  $k^{\text{th}}$  record
2. Temporarily remove  $(x_k, y_k)$  from the dataset
3. Train on the remaining  $(R-1)$  datapoints
4. Note your error  $(x_k, y_k)$

# LOOCV (Leave-one-out Cross Validation)

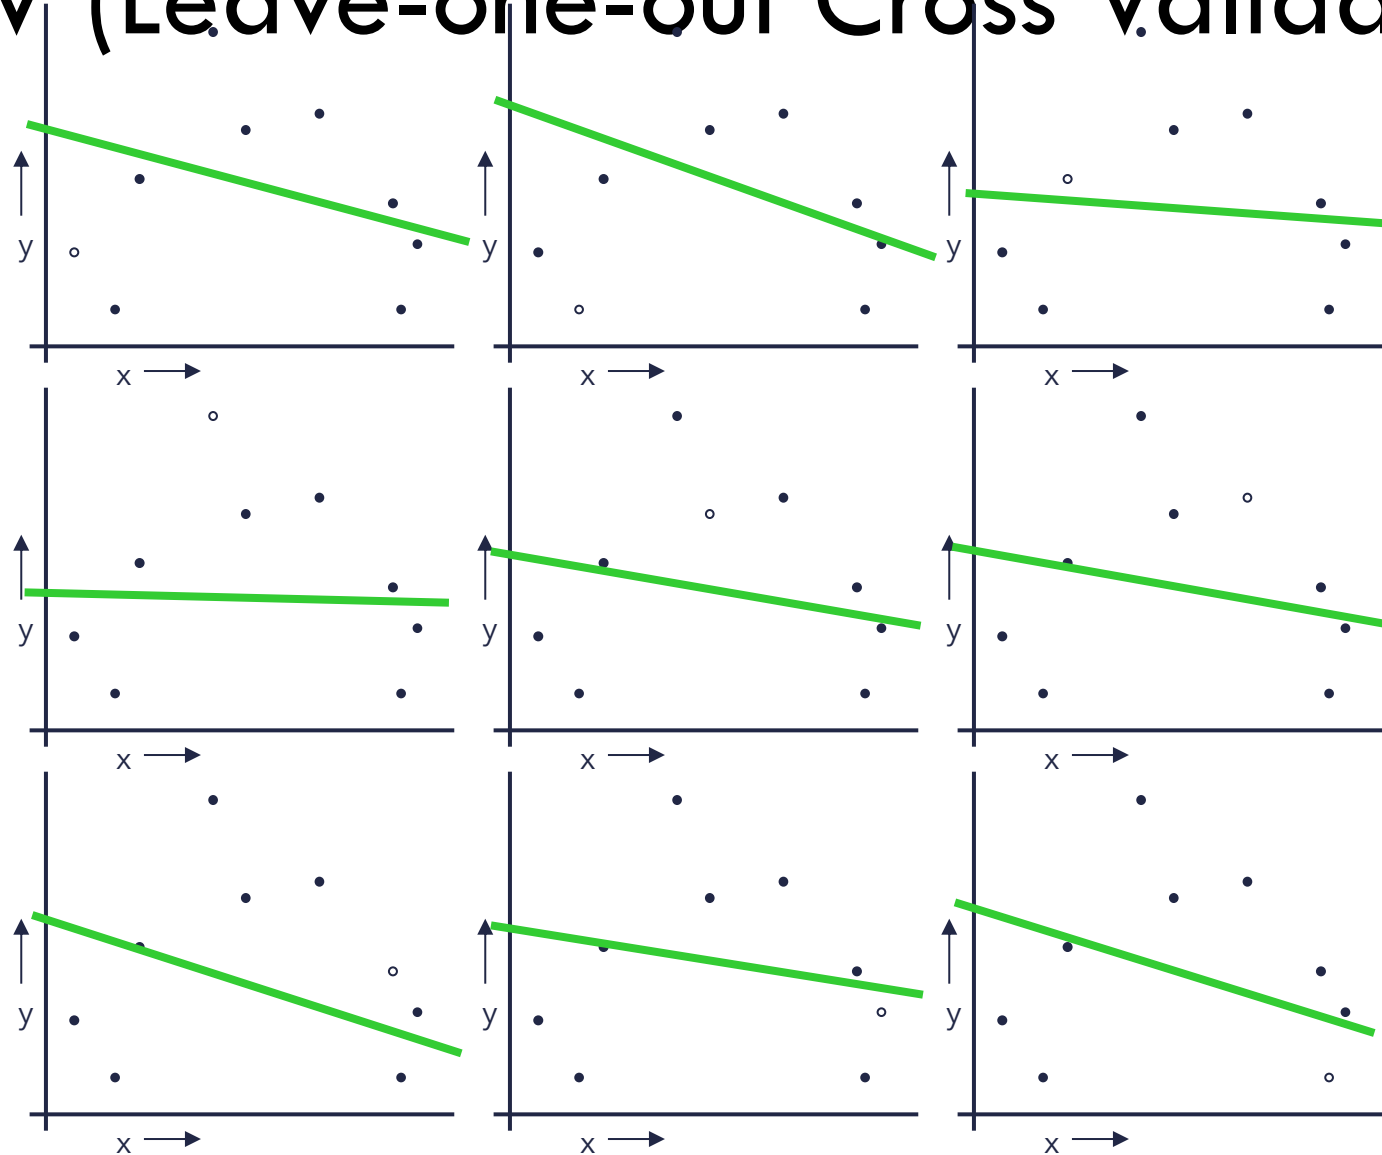


For  $k=1$  to  $R$

1. Let  $(x_k, y_k)$  be the  $k^{\text{th}}$  record
2. Temporarily remove  $(x_k, y_k)$  from the dataset
3. Train on the remaining  $(R-1)$  datapoints
4. Note your error  $(x_k, y_k)$

When you've done all points, report the mean error

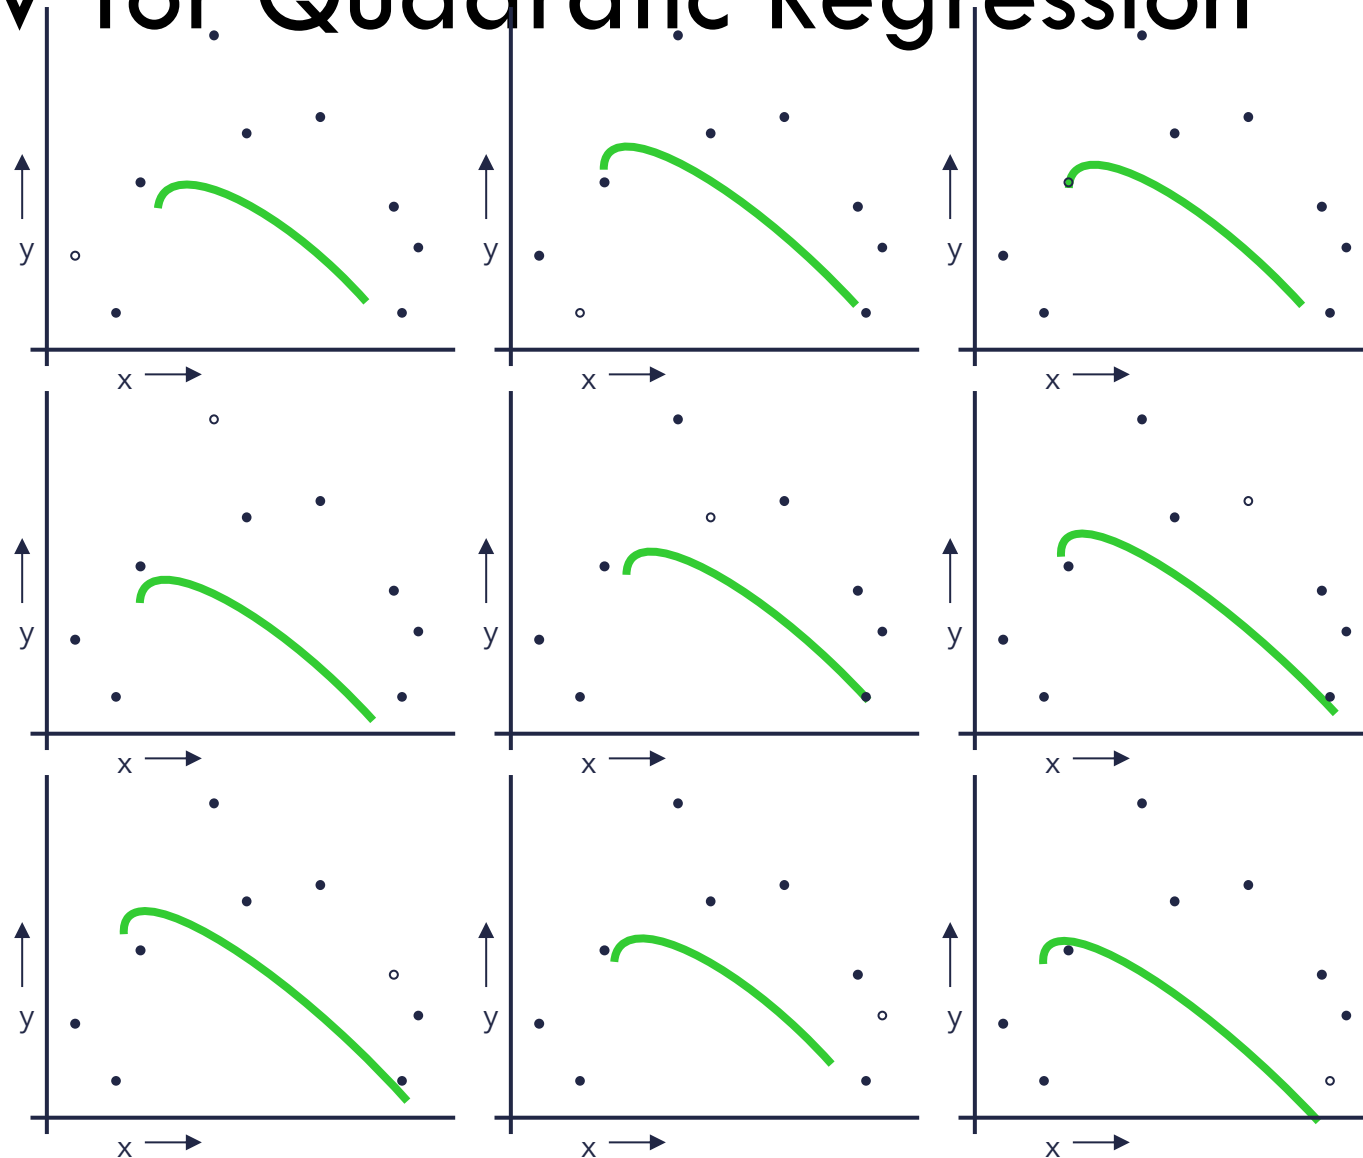
# LOOCV (Leave-one-out Cross Validation)



$$MSE_{LOOCV} = 2.12$$

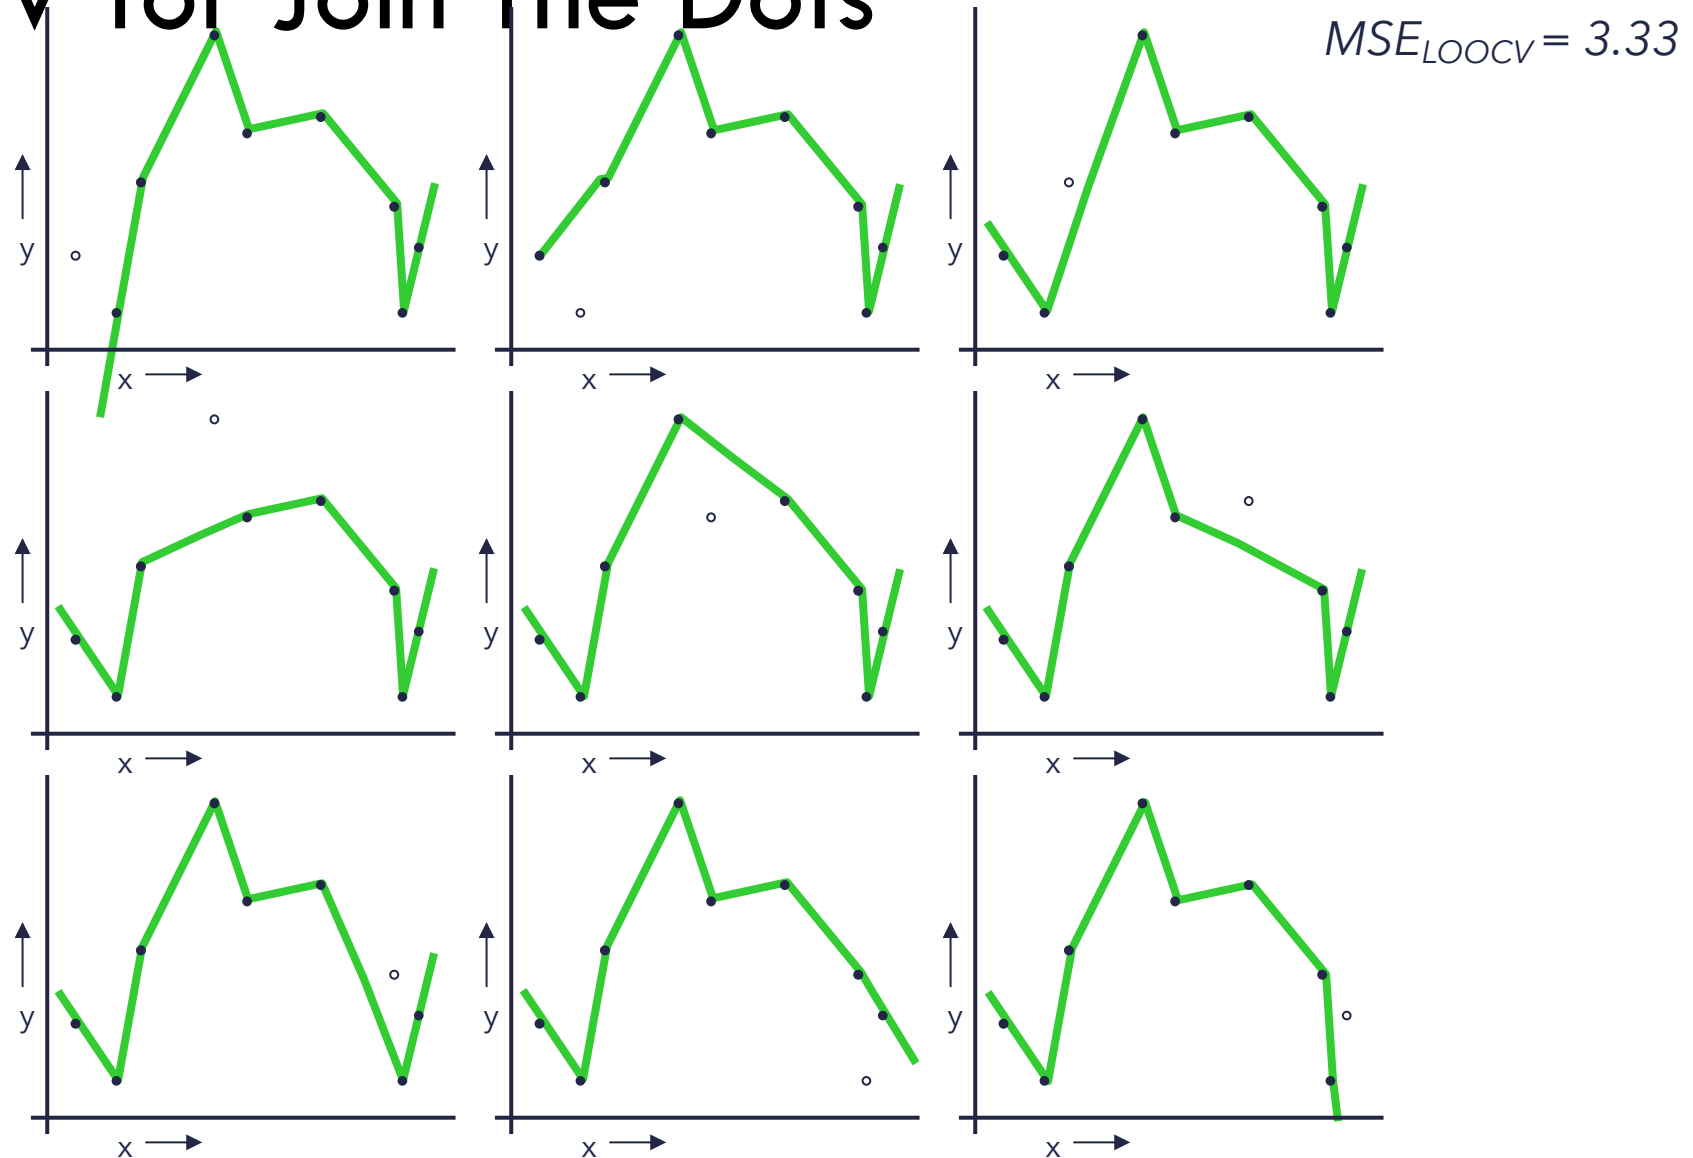
$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

# LOOCV for Quadratic Regression



$$MSE_{LOOCV} = 0.962$$

# LOOCV for Join The Dots

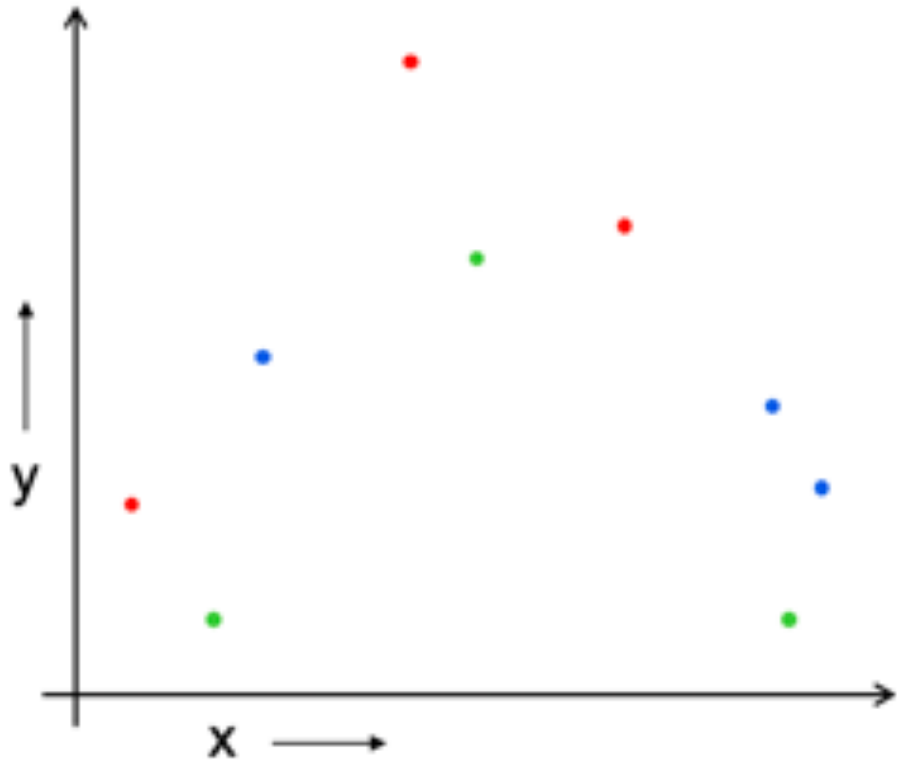


# Which kind of Validation?

	Downside	Upside
<b>Test-set 70%- 30%</b>	Variance: unreliable estimate of future performance	Cheap
<b>Leave- one-out</b>	Expensive	Doesn't waste data

...can we get the best of both worlds?

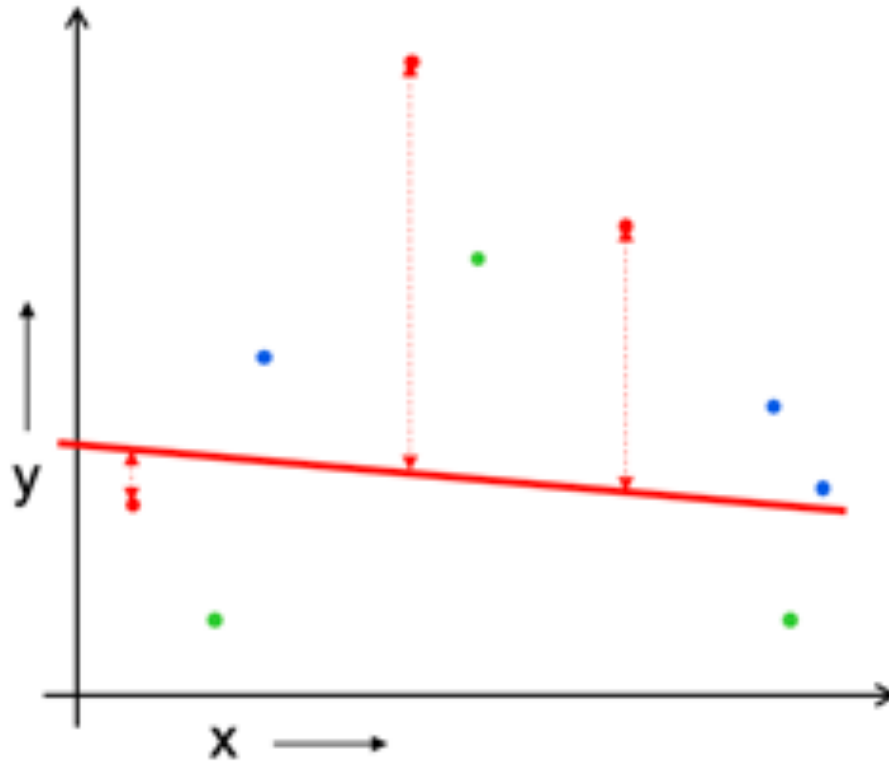
# k-fold Cross Validation



Randomly break the dataset into  $k$  partitions (in our example we'll have  $k=3$  partitions colored Red, Green and Blue)



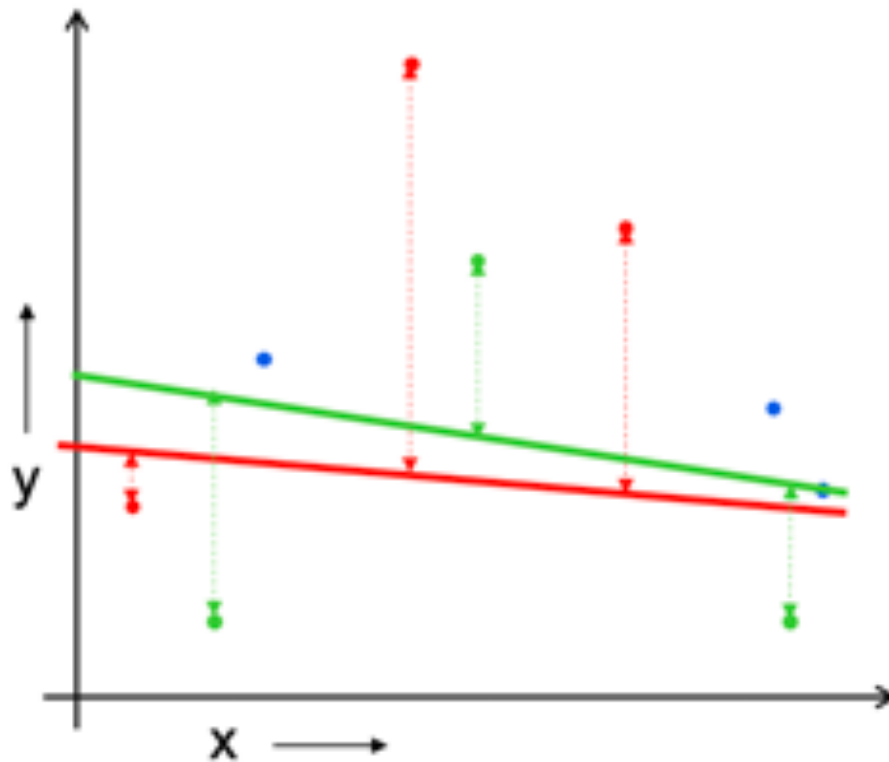
# k-fold Cross Validation



Randomly break the dataset into  $k$  partitions (in our example we'll have  $k=3$  partitions colored Red, Green and Blue)

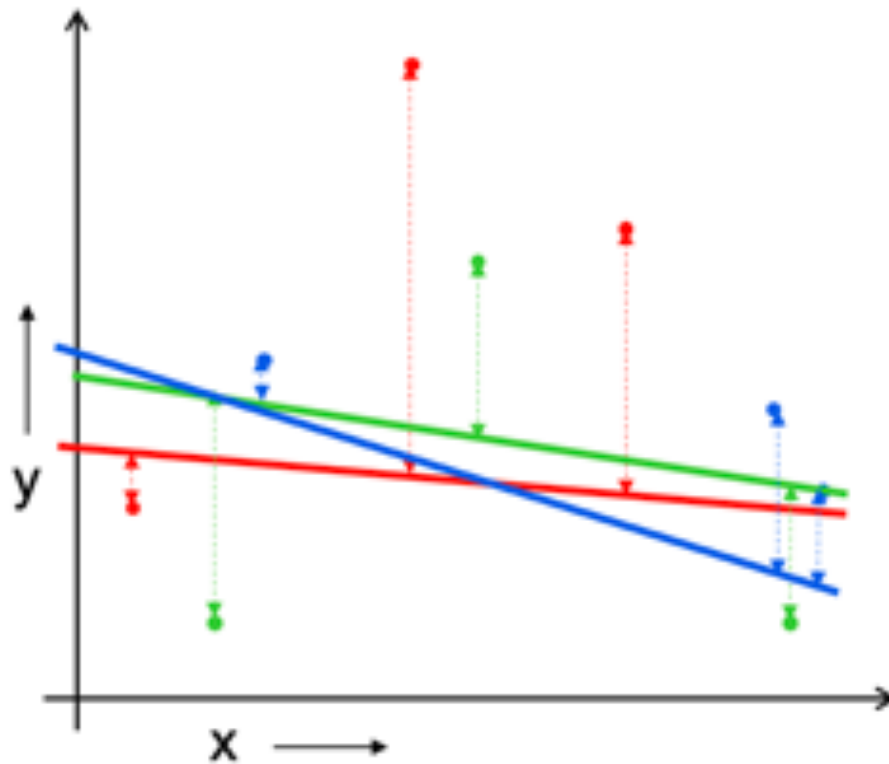
For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

# k-fold Cross Validation



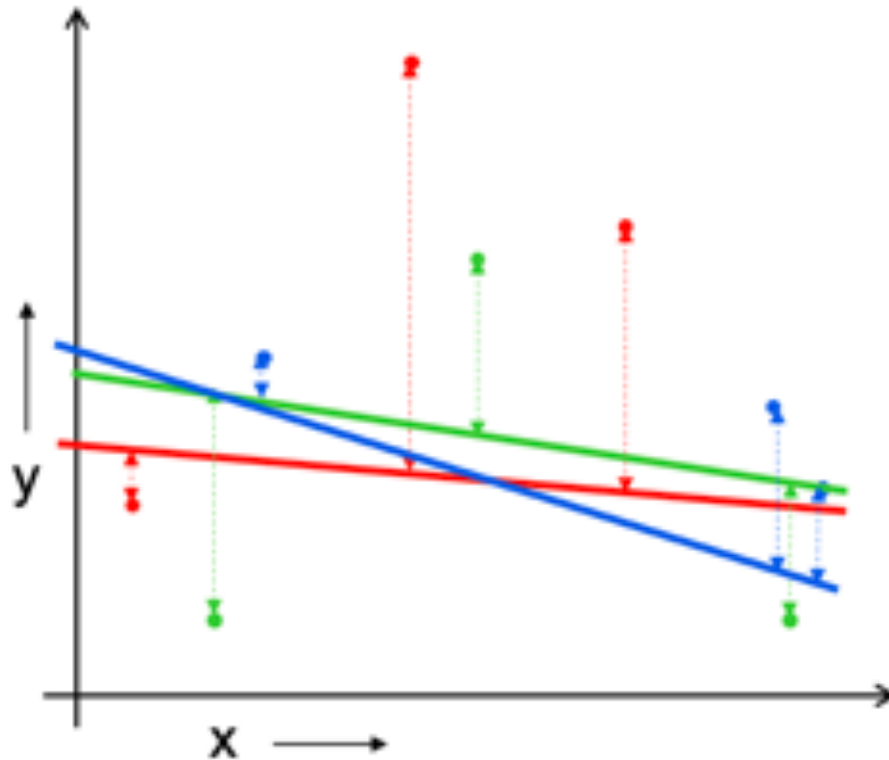
Randomly break the dataset into  $k$  partitions (in our example we'll have  $k=3$  partitions colored Red, Green and Blue)  
For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

# k-fold Cross Validation



Randomly break the dataset into  $k$  partitions (in our example we'll have  $k=3$  partitions colored **Red**, **Green** and **Blue**)  
For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.

# k-fold Cross Validation

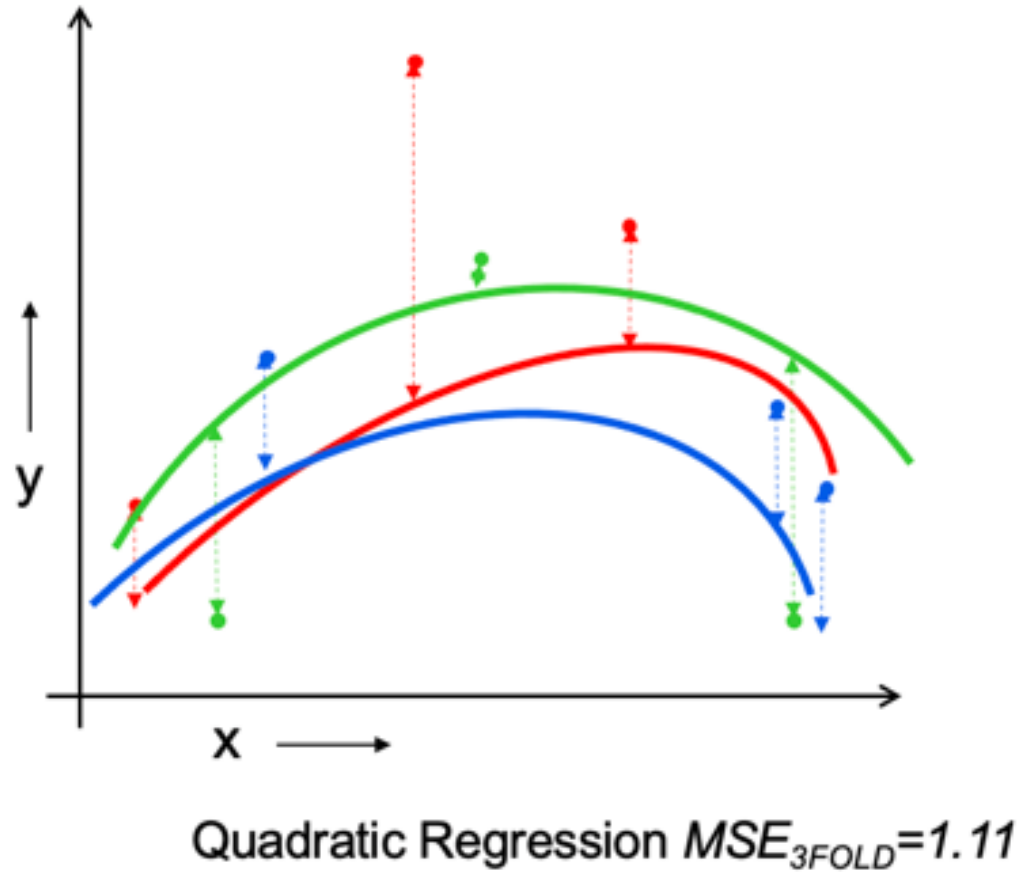


Randomly break the dataset into  $k$  partitions (in our example we'll have  $k=3$  partitions colored Red, Green and Blue)

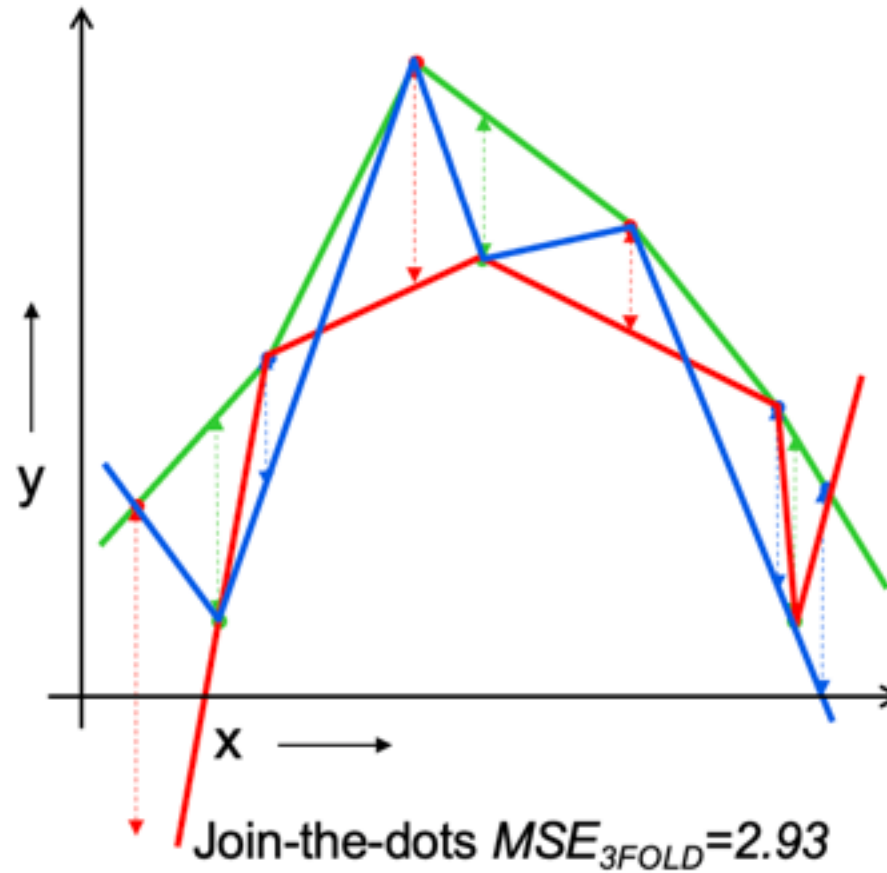
Then report the mean error.

Linear Regression  $MSE_{3FOLD}=2.05$

# k-fold Cross Validation



# k-fold Cross Validation

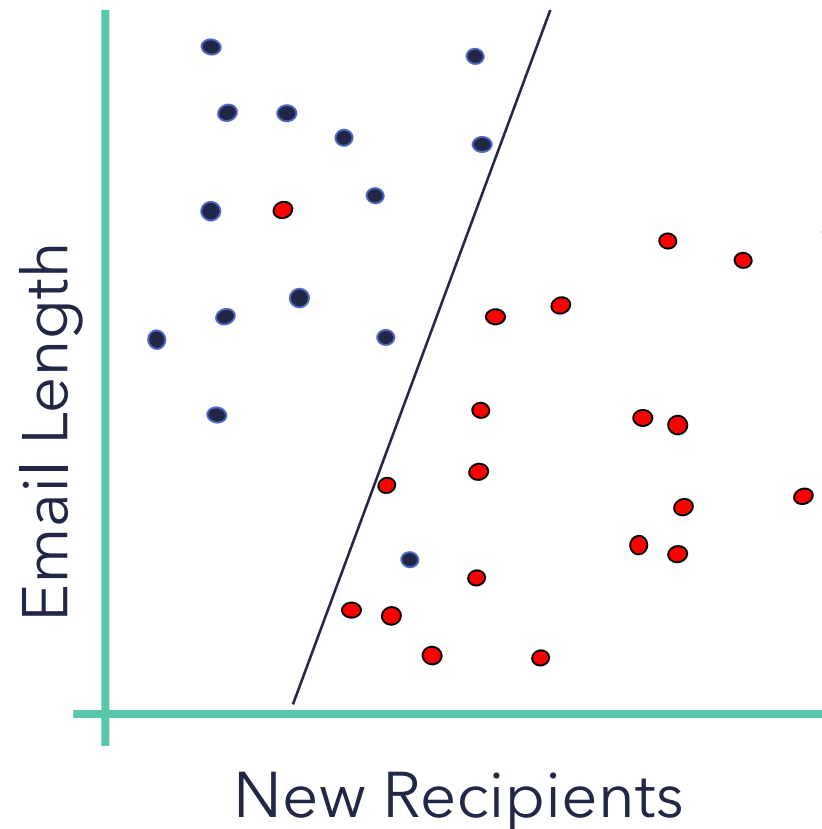


# Which kind of Cross Validation?

	Downside	Upside
<b>Test-set</b>	Variance: unreliable estimate of future performance.	Cheap.
<b>Leave-one-out</b>	Expensive. Has some weird behavior.	Doesn't waste data.
<b>10-fold</b>	Wastes 10% of the data. 10 times more expensive than test set.	Only wastes 10%. Only 10 times more expensive instead of R times.
<b>3-fold</b>	Wastes more data than 10-fold. More Expensive than test set.	Slightly better than test-set.
<b>R-fold</b>	Identical to Leave-one-out.	

# Evaluating what has been learned

1. We randomly select a portion of the data to be used for training (the training set)
2. Train the model on the training set
3. Once the model is trained, we run the model on the remaining instances (the test set) to see how it performs



**Confusion Matrix**

Classified as

	Black	Red
Black	7	1
Red	0	5



# Not by Accuracy alone

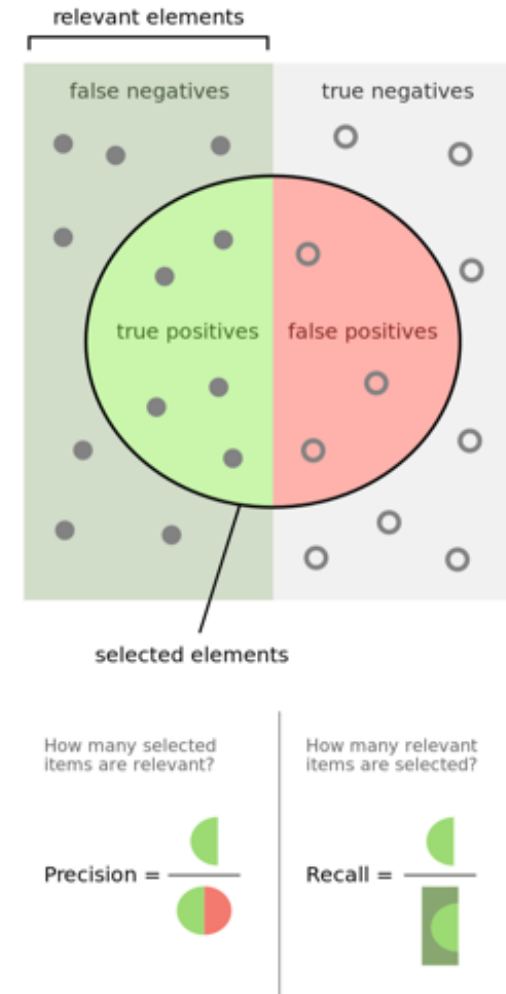
- Higher accuracy does not necessarily simply better performance on target task
- Classes are often very skewed
  - Malicious insider is rare
  - What if 99:1 split?
  - Always say "no" -- 99% default accuracy
- Errors have different costs
  - Deleting an email vs. machine infected

# Other

True Class	Predicted class	
	Yes	No
Yes	TP: True Positive	FN: False Negative
No	FP: False Positive	TN: True Negative

- Error rate = # of errors / # of instances =  $(FN+FP) / N$
- Recall (TPR) = # of found positives / # of positives
  - =  $TP / (TP+FN)$  = sensitivity = hit rate
- Precision = # of found positives / # of found =  $TP / (TP+FP)$
- Specificity (TNR) =  $TN / (TN+FP)$
- False alarm rate (FPR) =  $FP / (FP+TN)$  = 1 - Specificity

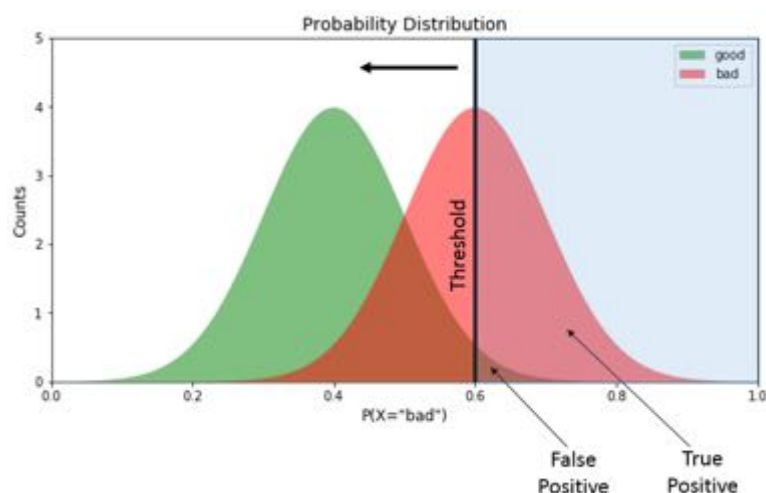
# Precision & Recall



Taken from: [https://en.wikipedia.org/wiki/Precision\\_and\\_recall](https://en.wikipedia.org/wiki/Precision_and_recall)

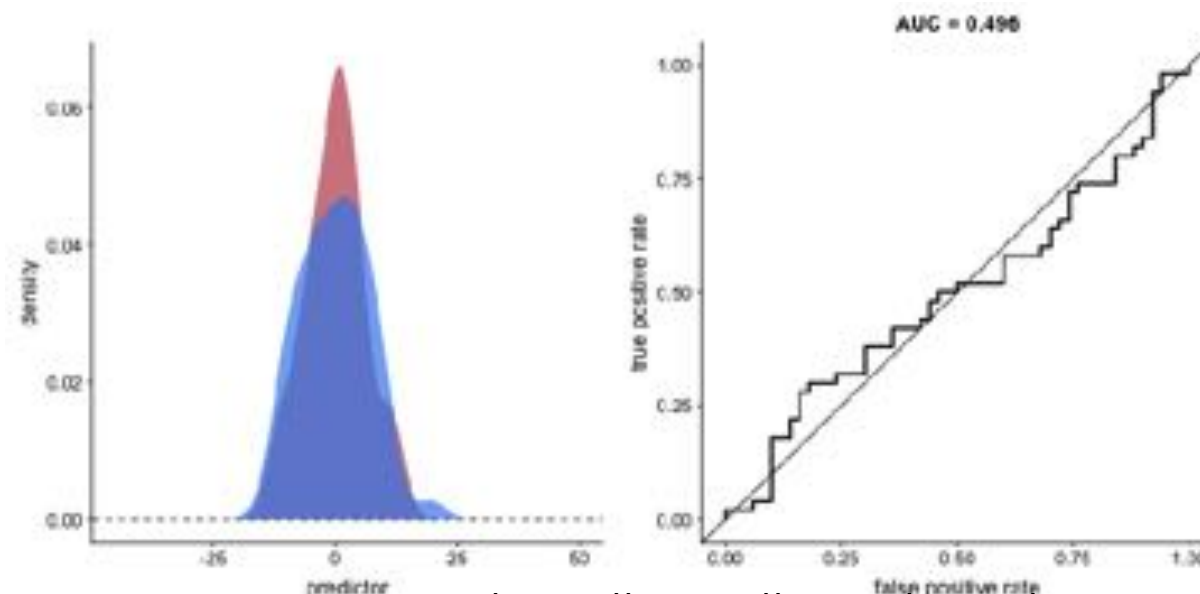
# ROC Curve

- Receiver Operating Characteristic (ROC) curve:
  - A technique for visualizing, organizing and selecting classifiers based on their performance
  - Two-dimensional graph in which the TPR (Recall) is plotted on the Y axis and the FPR (False alarm rate) is plotted on the X axis
  - Depicts relative tradeoffs between benefits (true positives) and costs (false positives)

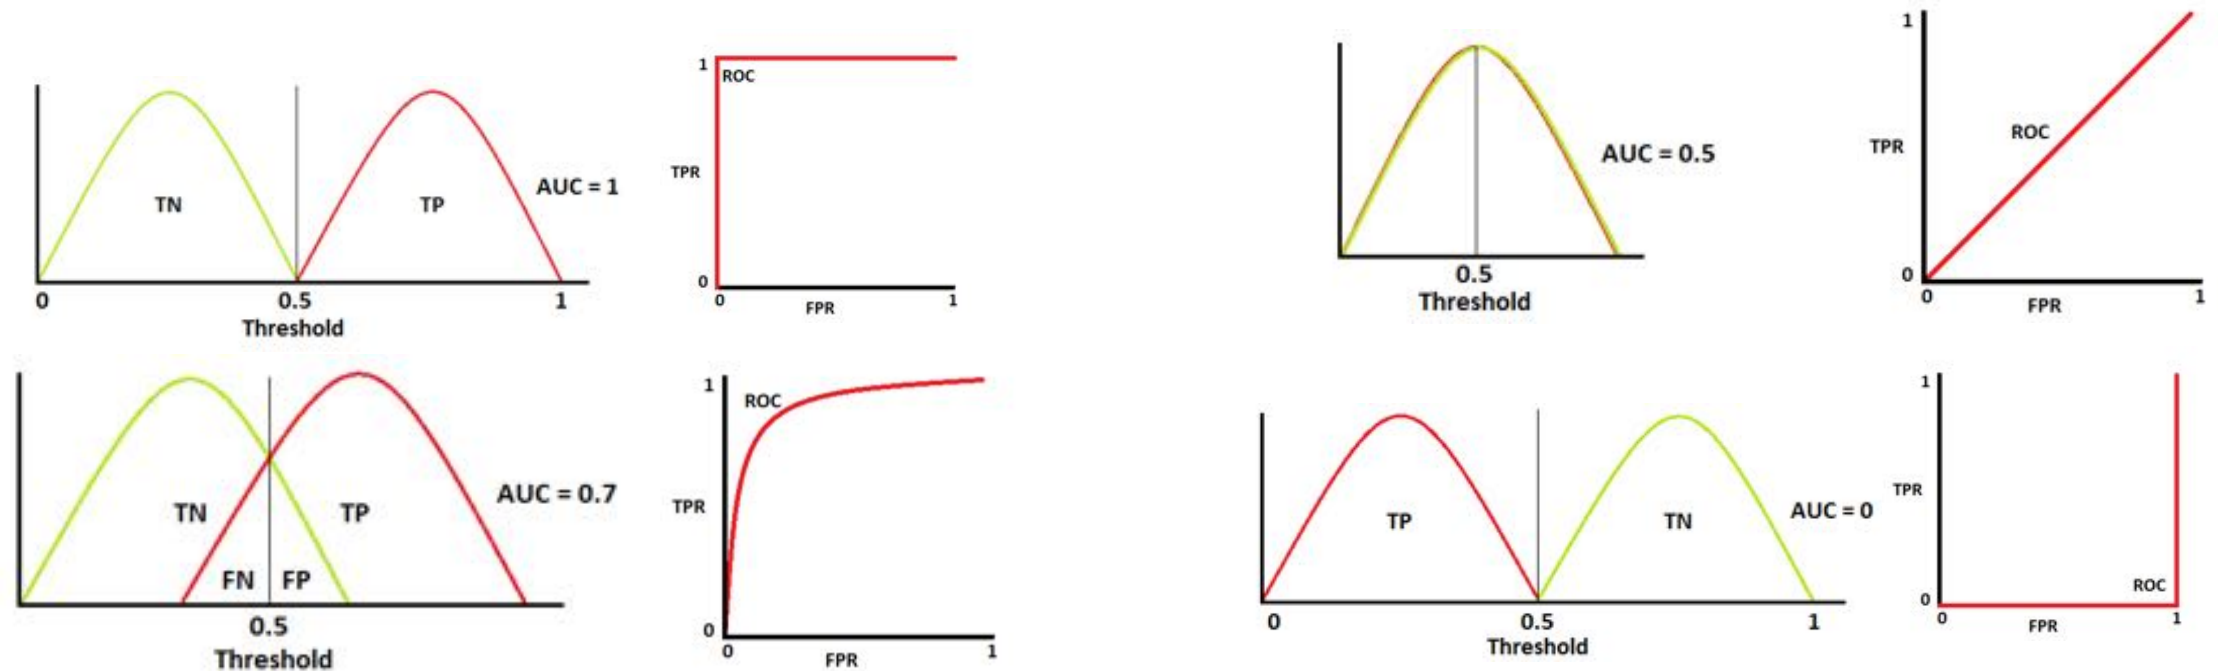


$$TPR = \frac{\text{True Positive}}{\text{Total Positive}}$$

$$FPR = \frac{\text{False Positive}}{\text{Total Negative}}$$

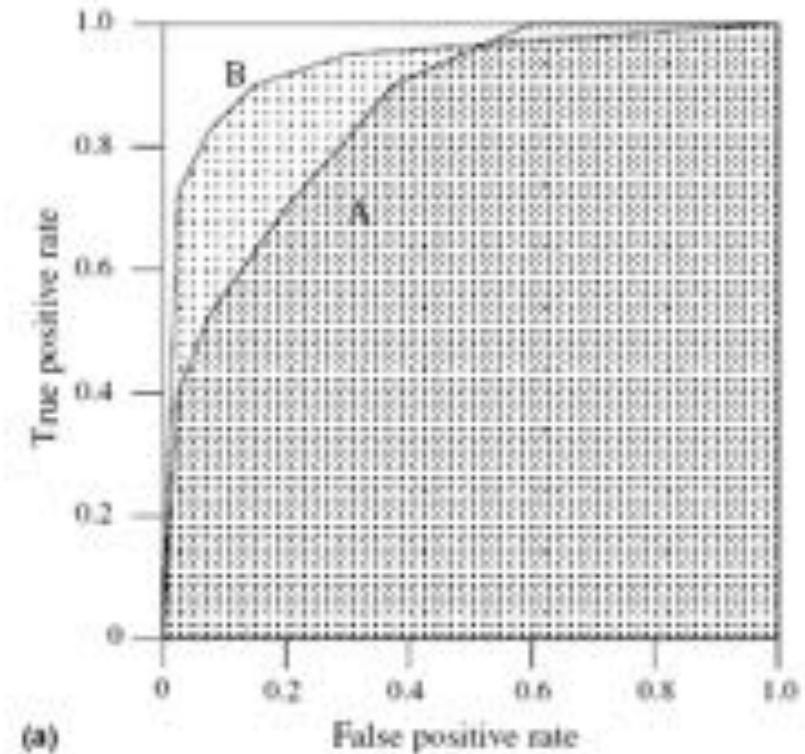


# ROC Curve For Binary Classification



# Area under ROC curve (AUC)

- Comparing two ROC curves:
- The graph represents the areas under two ROC curves, A and B
- Classifier B has greater area and therefore better average performance



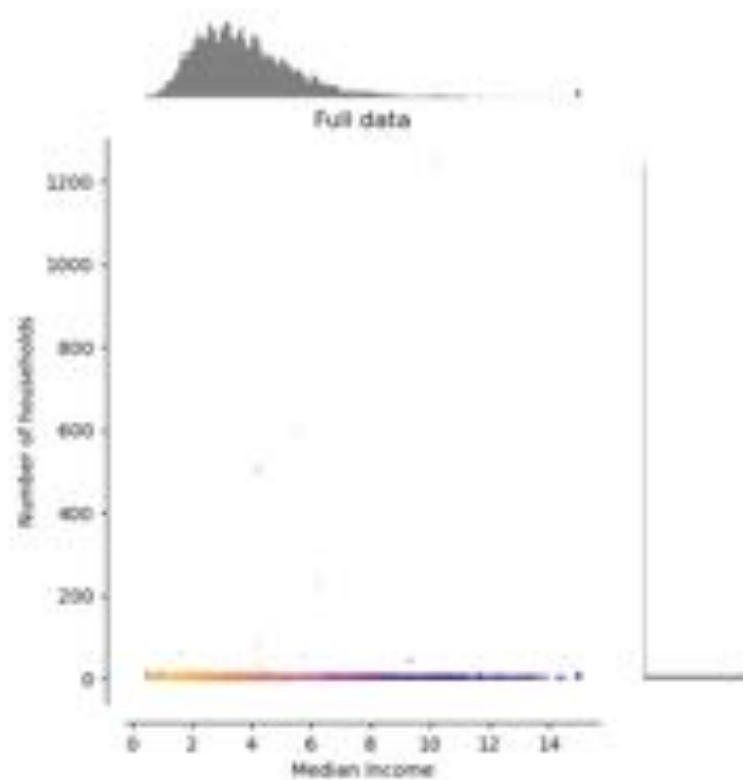


# Normalization

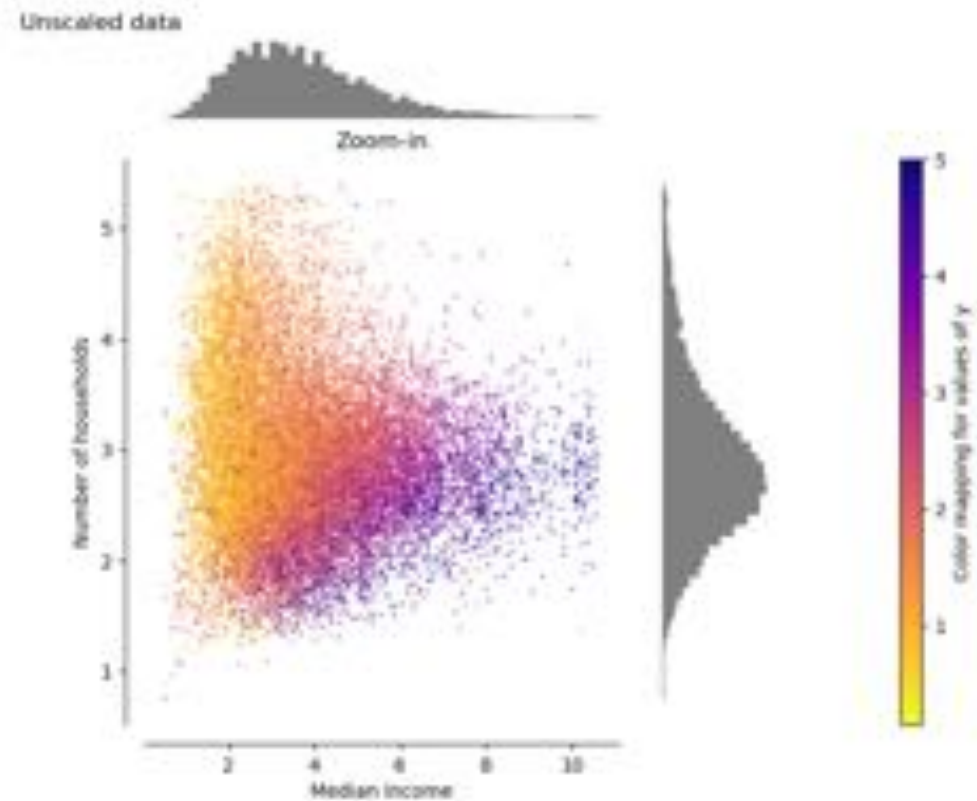
Please also check the IPYTHON

# Normalization

Original Data



Zoom in without outliers



Source: [https://scikit-learn.org/stable/auto\\_examples/preprocessing/plot\\_all\\_scaling.html#sphx-glr-auto-examples-preprocessing-plot-all-scaling-py](https://scikit-learn.org/stable/auto_examples/preprocessing/plot_all_scaling.html#sphx-glr-auto-examples-preprocessing-plot-all-scaling-py)



# StandardScaler

- Removes the mean and scales the data to unit variance.
  - Unit variance means that the standard deviation of a sample as well as the variance will tend towards 1 as the sample size tends towards infinity.
  - This is done per dimension --> each dim will have average  $\mu = 0$  and  $\sigma = 1$ .
- The scaling shrinks the range of the feature values as shown in the left figure below. However, the outliers have an influence when computing the empirical mean and standard deviation.
- [StandardScaler](#) therefore cannot guarantee balanced feature scales in the presence of outliers.

Standardization:

$$z = \frac{x - \mu}{\sigma}$$

with mean:

$$\mu = \frac{1}{N} \sum_{i=1}^N (x_i)$$

and standard deviation:

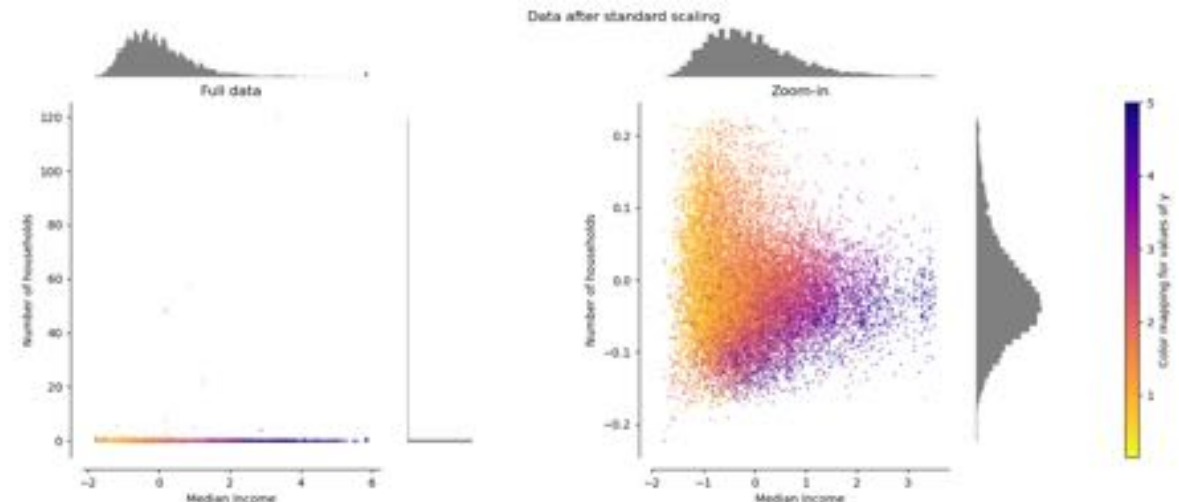
$$\sigma = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2}$$

```
from sklearn.preprocessing import StandardScaler
import numpy as np

# 4 samples/observations and 2 variables/features
data = np.array([[0, 0], [1, 0], [0, 1], [1, 1]])
scaler = StandardScaler()
scaled_data = scaler.fit_transform(data)

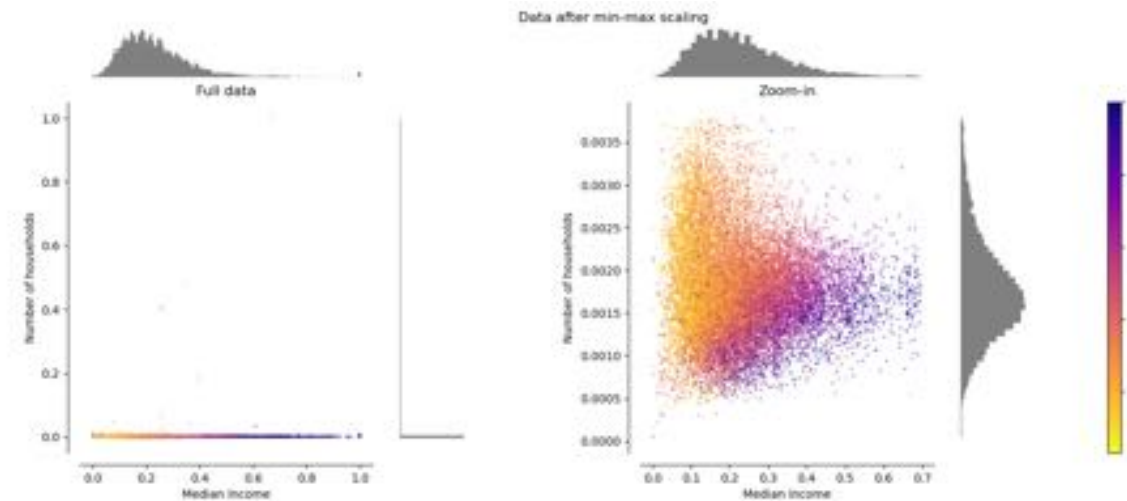
print(data)
[[0, 0],
 [1, 0],
 [0, 1],
 [1, 1]]

print(scaled_data)
[[-1. -1.]
 [ 1. -1.]
 [-1.  1.]
 [ 1.  1.]
```



# MinMaxScaler

- rescales the data set such that all feature values are in the range [0, 1] as shown in the right panel below.
- However, this scaling compresses all inliers into the narrow range [0, 0.005] for the transformed number of households.
- Both [StandardScaler](#) and [MinMaxScaler](#) are very sensitive to the presence of outliers.
- $X\_std = (X - X.min(axis=0)) / (X.max(axis=0) - X.min(axis=0))$   
 $X\_scaled = X\_std * (max - min) + min$





# Feature Selection

# Why feature selection?

- A large number of extracted features, some of which redundant or irrelevant, present several problems
- Misleading the learning algorithm
- Over-fitting and therefore reducing generality
- Increasing model complexity and processing time
- Large number of features requires more examples
- Requires more resources

# Goal

- Select a subset of relevant features that eventually will lead to a better, faster, and easier to understand model

Number of new Recipients	Email Length (K)	Country (IP)	Customer Type	Email Type
0	2	Germany	Gold	Ham
1	4	Germany	Silver	Ham
5	2	Nigeria	Bronze	Spam
2	4	Russia	Bronze	Spam
3	4	Germany	Bronze	Ham
0	1	USA	Silver	Ham
4	2	USA	Silver	Spam

# Why feature selection?

- Assume  $m$  features
- We want to select  $k$

$$\binom{m}{k} = \frac{m!}{k! (m - k)!}$$

- Example:  $m=20$   $k=5 \rightarrow 15,504$  combinations

# Feature selection methods

- Univariate vs. Multivariate
- Univariate - selecting the best features based on univariate statistical tests.
  - It can be seen as a preprocessing step to an estimator.

# Univariate method

- Assign a rank to each feature independently
- Indicates how good the feature is in discriminating between classes
- Differ in the ranking metric
- Remove useless



# Low Variance

- Simple baseline approach to feature selection
- It **removes** all **features** whose **variance** doesn't meet some **threshold**
- It removes all zero-variance features, i.e. features that have the same value in all samples.
- **Example:** remove features that are 1 or 0 in 80% of the samples.
  - **Boolean** features are **Bernoulli random** variables, and the variance of such variables is given by:

$$\text{Var}[X] = p(1 - p)$$

- $\text{Thr} = 0.8 * (1 - 0.8)$

```
>>> from sklearn.feature_selection import VarianceThreshold
>>> X = [[0, 0, 1], [0, 1, 0], [1, 0, 0], [0, 1, 1], [0, 1, 0], [0, 1, 1]]
>>> sel = VarianceThreshold(threshold=(.8 * (1 - .8)))
>>> sel.fit_transform(X)
array([[0, 1],
       [1, 0],
       [0, 0],
       [1, 1],
       [1, 0],
       [1, 1]])
```

- As expected the VarianceThreshold has removed the first column which has probability  $p = 5/6 > 0.8$  of containing zeros

# SelectKBest

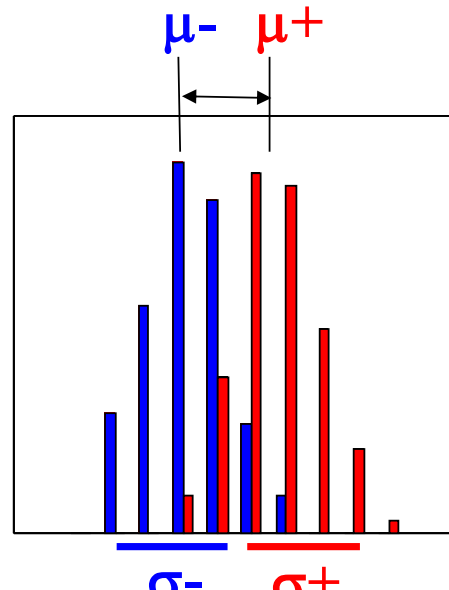
- Removes all but the K highest scoring features
- Select based on different functions
  - Mutual information
  - Chi2
  - Note different functions can work better for specific use cases
- Return univariate scores and p-values (for most variants)

```
>>> from sklearn.datasets import load_iris
>>> from sklearn.feature_selection import SelectKBest
>>> from sklearn.feature_selection import chi2
>>> X, y = load_iris(return_X_y=True)
>>> X.shape
(150, 4)
>>> X_new = SelectKBest(chi2, k=2).fit_transform(X, y)
>>> X_new.shape
(150, 2)
```

# Univariate method

## Fisher Score

- **Calculates** the difference, described in terms of **mean** and **standard deviation**, between the **positive and negative examples** relative to a certain feature.
- **Features with high quality** should assign **similar values** to instances in the **same class** and **different values** to instances from **different classes**.



$$R^{(i)} = \frac{\left| \mu_{+}^{(i)} - \mu_{-}^{(i)} \right|}{\sigma_{+}^{(i)} + \sigma_{-}^{(i)}}$$

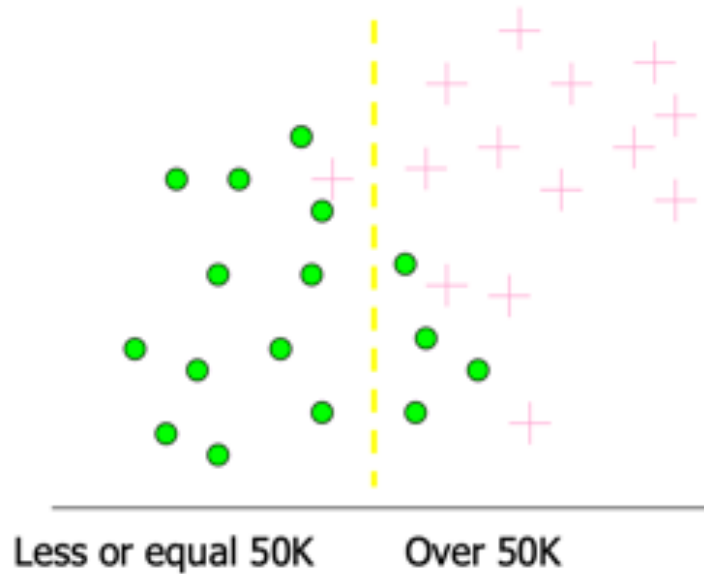
# Entropy & Information Gain Explained

- Information Gain & Entropy slides taken from Prof. Linda Shapiro
- Source:  
<https://homes.cs.washington.edu/~shapiro/EE596/notes/InfoGain.pdf>

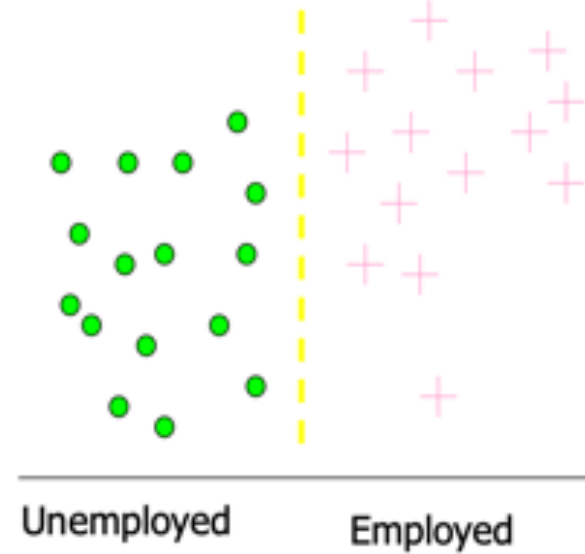
# Entropy

Which test is more informative?

**Split over whether  
Balance exceeds 50K**



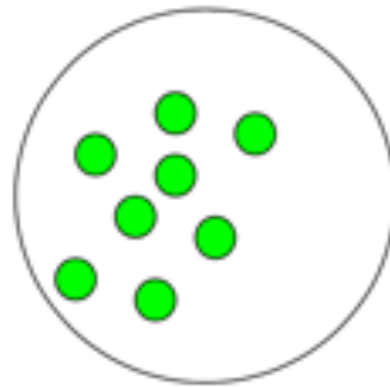
**Split over whether  
applicant is employed**



# Entropy

## Impurity/Entropy (informal)

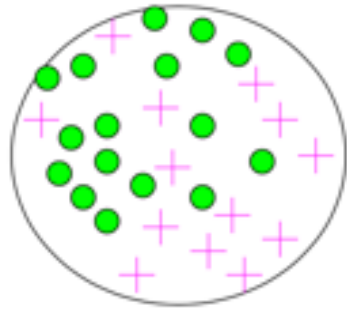
- Measures the level of **impurity** in a group of examples



# Entropy

## Impurity

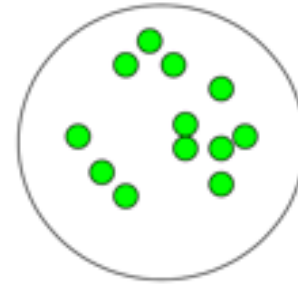
**Very impure group**



**Less impure**



**Minimum  
impurity**



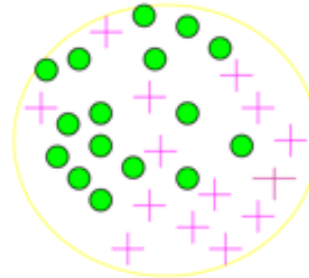
# Entropy

$$\text{Log}_2 1 = 0$$

- Entropy =  $\sum_i -p_i \log_2 p_i$

$p_i$  is the probability of class  $i$

Compute it as the proportion of class  $i$  in the set.



16/30 are green circles; 14/30 are pink crosses

$$\log_2(16/30) = -.9; \quad \log_2(14/30) = -1.1$$

$$\text{Entropy} = -(16/30)(-.9) - (14/30)(-1.1) = .99$$

- Entropy comes from information theory. The higher the entropy the more the information content.

What does that mean for learning from examples?



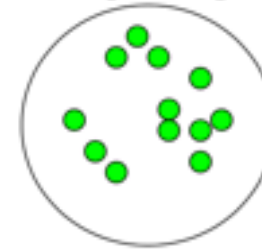
# Entropy – 2 class cases

- What is the entropy of a group in which all examples belong to the same class?

– entropy =  $-1 \log_2 1 = 0$

not a good training set for learning

**Minimum  
impurity**



- What is the entropy of a group with 50% in either class?

– entropy =  $-0.5 \log_2 0.5 - 0.5 \log_2 0.5 = 1$

good training set for learning

**Maximum  
impurity**

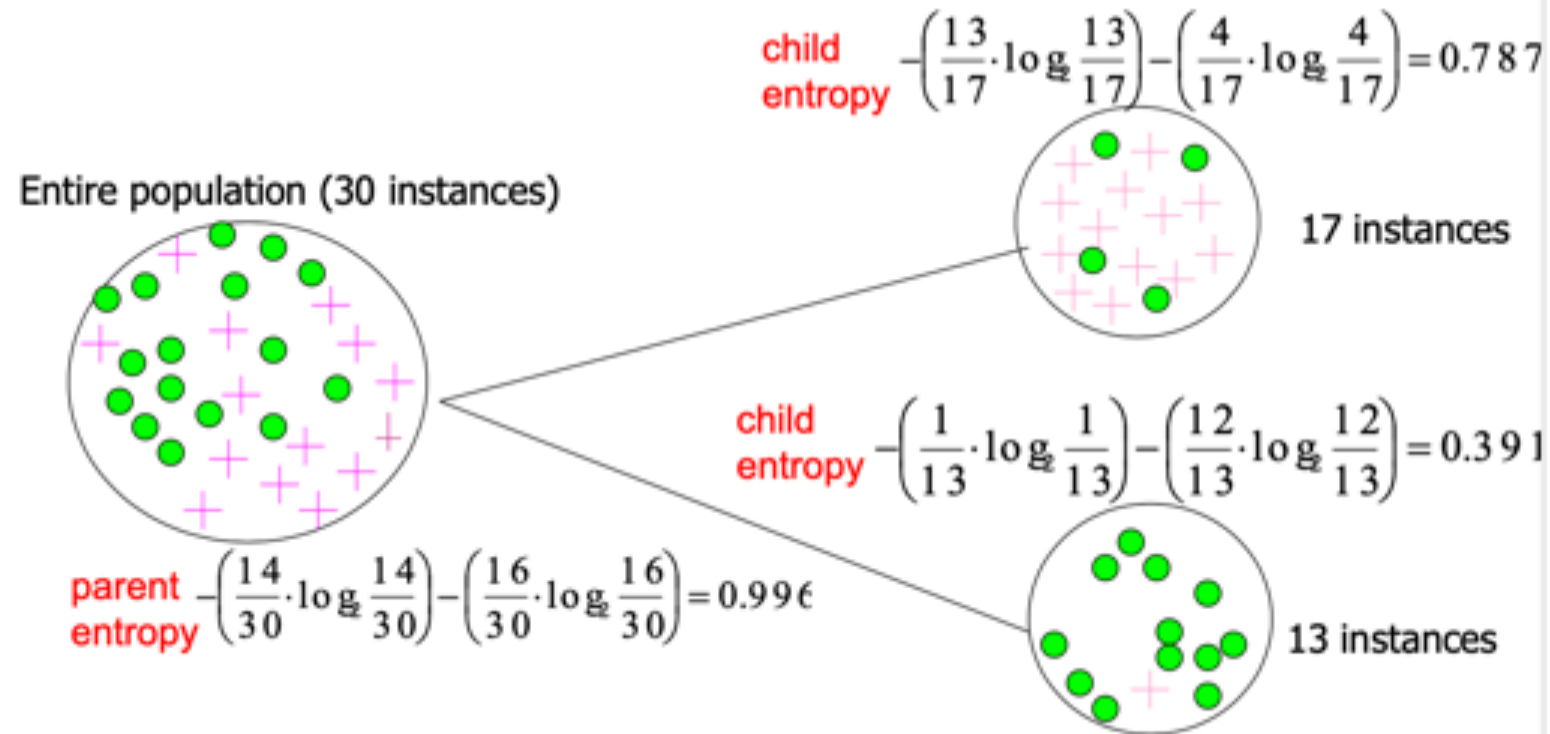


# Information Gain

- We want to determine **which attribute** in a given set of training feature vectors is **most useful** for discriminating between the classes to be learned.
- **Information gain** tells us how important a given attribute of the feature vectors is.
- We will use it to decide the ordering of attributes in the nodes of a decision tree.

# Information Gain

**Information Gain** = entropy(parent) – [average entropy(children)]



**(Weighted) Average Entropy of Children** =  $\left(\frac{17}{30} \cdot 0.787\right) + \left(\frac{13}{30} \cdot 0.391\right) = 0.615$

**Information Gain = 0.996 - 0.615 = 0.38 for this split**

# Information Gain Example

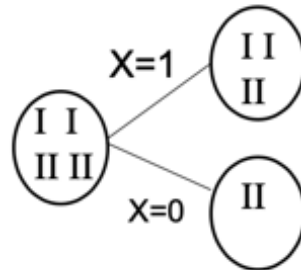
Training Set: 3 features and 2 classes

X	Y	Z	C
1	1	1	I
1	1	0	I
0	0	1	II
1	0	0	II

# Information Gain Example

X	Y	Z	C
1	1	1	I
1	1	0	I
0	0	1	II
1	0	0	II

Split on attribute X



If X is the best attribute,  
this node would be further split.

$$\begin{aligned} E_{\text{child1}} &= -(1/3)\log_2(1/3) - (2/3)\log_2(2/3) \\ &= .5284 + .39 \\ &= .9184 \end{aligned}$$

$$E_{\text{child2}} = 0$$

$$E_{\text{parent}} = 1$$

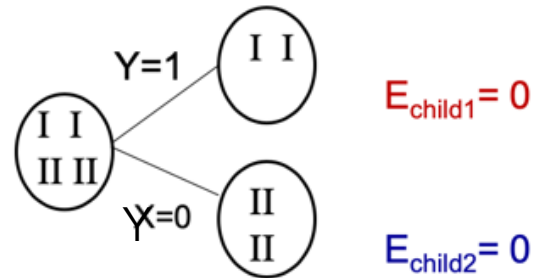
$$\text{GAIN} = 1 - (3/4)(.9184) - (1/4)(0) = .3112$$

# Information Gain Example

X	Y	Z	C
1	1	1	I
1	1	0	I
0	0	1	II
1	0	0	II

Note: The entropy here is 0 but Information Gain Is 1

Split on attribute Y



$E_{\text{child1}} = 0$

$E_{\text{child2}} = 0$

$E_{\text{parent}} = 1$

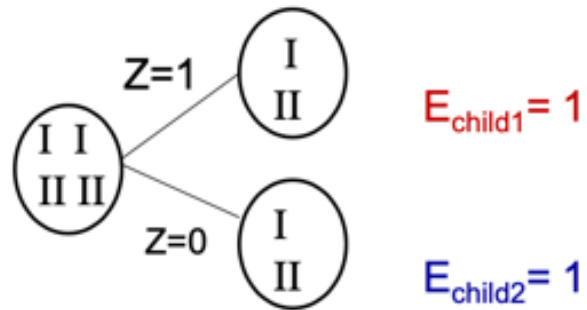
$\text{GAIN} = 1 - (1/2)0 - (1/2)0 = 1$ ; BEST ONE

# Information Gain Example

X	Y	Z	C
1	1	1	I
1	1	0	I
0	0	1	II
1	0	0	II

Note: Here the Entropy is 1  
But the Information Gain is 0

Split on attribute Z



$$E_{\text{parent}} = 1$$

$$\text{GAIN} = 1 - (1/2)(1) - (1/2)(1) = 0 \quad \text{ie. NO GAIN; WORST}$$